

System-Level Design (and Modeling for Embedded Systems)

Lecture 3 – Models of Computation

Kim Grüttner kim.gruettner@dlr.de

Henning Schlender henning.schlender@dlr.de

Jörg Walter joerg.walter@offis.de

System Evolution and Operation
German Aerospace Center (DLR)
&

Distributed Computation and Communication
OFFIS



© 2009 Andreas Gerstlauer
Electrical and Computer Engineering
University of Texas at Austin
gerstl@ece.utexas.edu

- **Models of Computation (MoCs)**
 - Introduction
 - Programming models
 - Concurrency models
- **Process-based MoCs**
 - Process networks
 - Dataflow
 - Process calculi
- **State-based MoCs**
 - Finite state machines
 - Hierarchical, concurrent state machines
 - Process state machines

- **Conceptual, abstract description of system behavior**
 - Classification based on underlying characteristics
 - Computation and communication
 - Define a class of design models
 - Well-defined, formal definition and semantics
 - Functionality (data) and order (time)
 - Formal analysis and reasoning
 - Various degrees of complexity and expressiveness
 - Decomposition into pieces and their relationship
 - Objects and composition rules
 - Data and control flow
- **Analyzability and expressiveness of behavioral models**
 - Inherent or analyzable properties (liveness, fairness, ...)
 - Restrictions (Turing completeness, halting problem)
 - Efficient implementation

- **Imperative programming models**
 - Statements that manipulate program state, control flow
 - Sequential programming languages [C, C++, Java, ...]
- **Declarative programming models**
 - Rules for data manipulation, data flow
 - Functional programming [Haskell, Lisp]
 - Logic programming [Prolog]
- **Von Neumann execution**
 - Implicit concurrency (dependency analysis)
 - Implicit ordering (time)
 - Implicit state (set of variables)
 - Implicit termination (Turing complete)

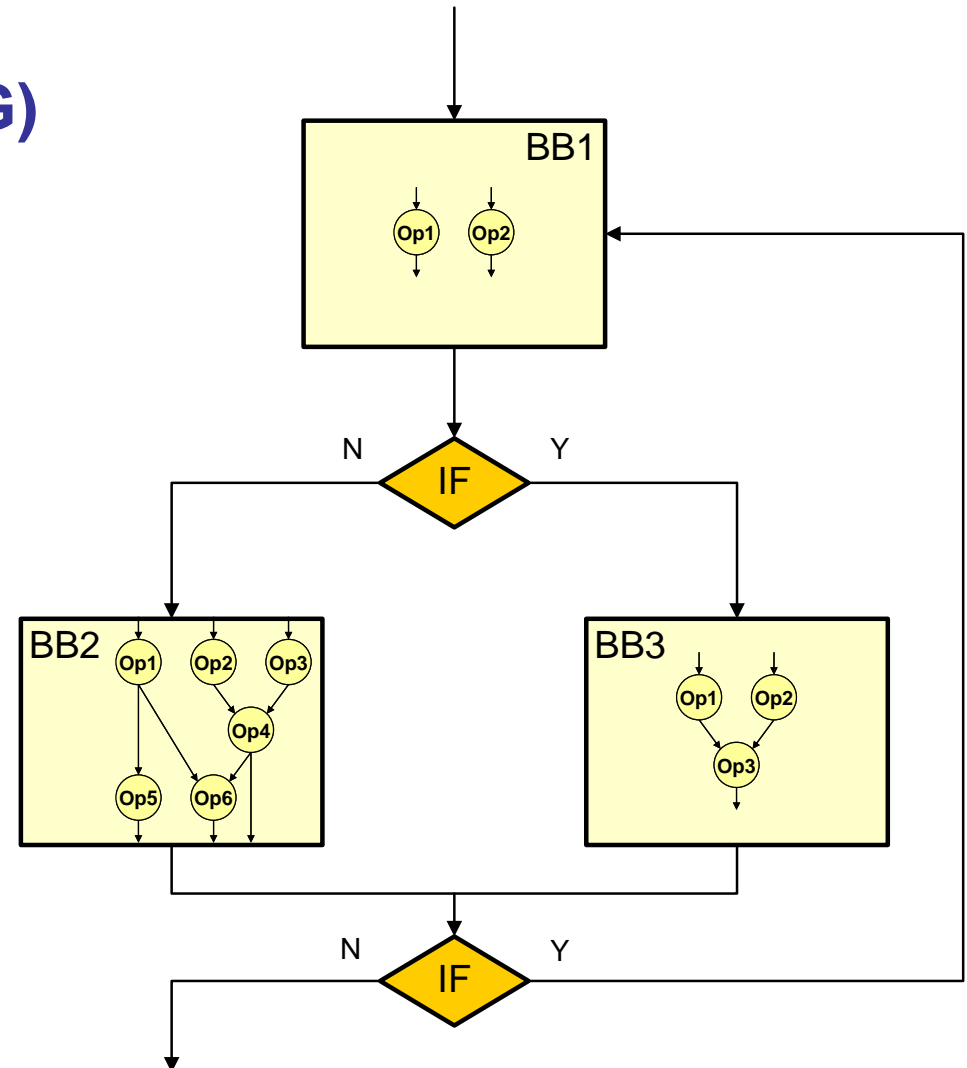
- **Control/Data Flow Graph (CDFG)**

- **Control**

- Branches, loops
 - Dependencies, order of basic blocks (BB) of computation

- **Data**

- Basic blocks (BBs) with expressions
 - Directed, acyclic graph of operations




➤ **Graphical representation (syntax) of program**

- **Transformative**



- Output = $F(\text{Input})$
 - Data processing

- **Reactive**




- Continuous interaction with environment
 - Control oriented: e.g. interactive (servers, GUIs) and constrained (embedded)

- **Ordering, interleaving of inputs and outputs (when)?**

- Timing defines reactive behavior (e.g. in case of feedback)

➤ Interaction, ordering relationship between blocks

- Control or data dependencies, communication (what)
- Synchronization, time (when)



➤ Physical: total order (real time)


- Distinct components naturally interleaved in (very fine) time

➤ Logical: partial order (logical time)

- Specification of interleaving based on causality only (implementation freedom)

- **Synchronous hypothesis**

- Reactions are instantaneous (zero time)
- Simultaneous (broadcast)
- Sequence of discrete steps (logical clock)
- Deterministic, static verifiable (independent of actual delays)



- **Challenges:**

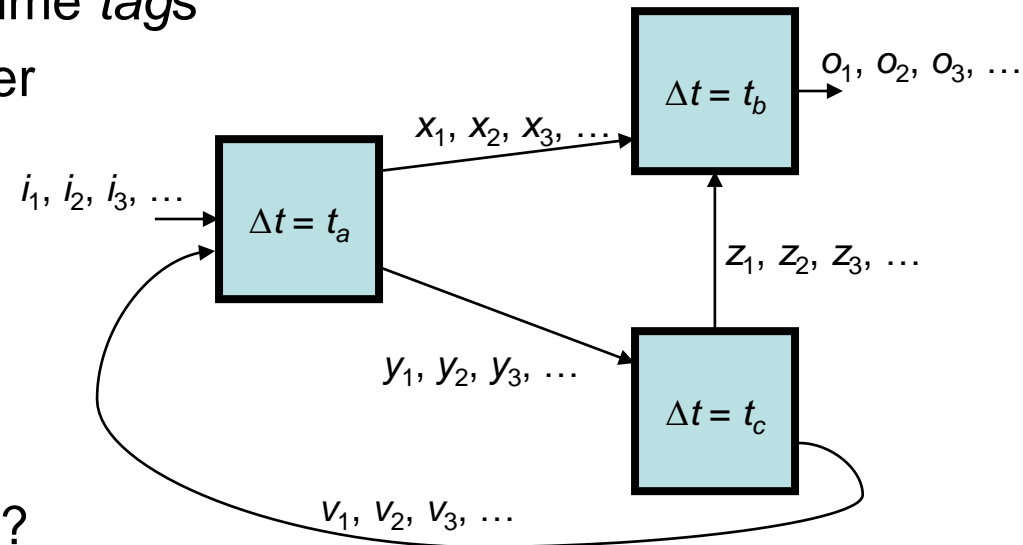
- Semantics: causal microsteps, but: cycles, conflicts?
- Implementation: $\Delta t \ll clock$, global clock, shared events

- **Synchronous languages**

- Imperative (control) [Esterel] or declarative (data) [Lustre] style
- Reject cycles [Lustre] or require unique fixed-point [Esterel]
- Hardware (synchronous) or software (safety critical) compilers

- **Asynchronous, relax relations**

- Signals = streams of events
- Event $e_i = (value, tag)$, discrete time $tags$
- Global event and evaluation order
- Flexible, arbitrary dynamic delays



- **Challenges**

- Semantics: simultaneous events?
- Implementation: maintain delays (global time)

- **Design languages**

- General, universal modeling and simulation (multi-scale)
- Hardware-description [VHDL, Verilog], system-level [SpecC, SystemC]
- Deterministic zero-delay delta steps, potentially endless cycles
- Ignore delays for synthesis, only realize causality (subset)

- **MoC examples**

- Programming models: imperative, declarative
 - Statements for algorithmic computation, operation-level granularity
- Concurrency models: synchronous, discrete event
 - Basic, general interaction and global relation between operations
- Process-based models: networks, dataflow, calculi
 - Activity, data flow/dependencies
- State-based models: evolution from FSM to PSM
 - State enumeration, control flow/dependencies

- **Specification and algorithm modeling**

- Ptolemy (UC Berkeley): heterogeneous mixture of MoCs
- Matlab/Simulink (Mathworks), LabView (NI): dataflow
- Statemate (IBM Rational), UML: StateCharts, HCFSM

✓ Models of Computation (MoCs)

- **Process-based MoCs**

- Process networks
- Dataflow
- Process calculi


- **State-based MoCs**

➤ Activity and causality (data flow)

- Asynchronous, coarse-grain concurrency


• Set of processes

- Processes execute in parallel
 - Concurrent composition
- Each process is internally sequential
 - Imperative program



• Inter-process communication


- Shared memory [Java]
 - Synchronization: critical section/mutex, monitor, ...
- Message passing [MPI]
 - Synchronous, rendezvous (blocking send)
 - Asynchronous, queues (non-blocking send)



➤ Implementation: OS processes or threads

- Single or multiple processors/cores

- **Circular chain of 2 or more processes which each hold a shared resource that the next one is waiting for**
 - Circular dependency through shared resources



- Prevent chain by using the same precedence
- Use timeouts (and retry), but: livelock
- **Dependency can be created when resources are shared**
 - Side effects, e.g. when blocking on filled queues/buffers

- **Deterministic: same inputs always produce same results**
- **Random: probability of certain behavior**
- **Non-deterministic: undefined behavior (for some inputs)**
 - Undefined execution order
 - Statement evaluation in imperative languages: $f(a++, a++)$
 - Concurrent process race conditions:

```
x = a;  
y = b;
```

```
a = 1;  
b = 2;
```

$x = ?, y = ?$

➤ Can be desired or undesired


➤ How to ensure correctness?

- Tedious and error-prone manual untangling (synchronization)
- Simulator will typically pick only one behavior

➤ But: over-specification?

- Leave freedom of implementation choice (concurrency)

- **C-like processes communicating via FIFO channels**
 - Unbounded, uni-directional, point-to-point queues
 - Sender (`send()`) never blocks
 - Receiver (`wait()`) blocks until data available



➤ **Deterministic**

- Behavior does not depend on scheduling strategy
- Focus on causality, not order (implementation independent)

- **Determinism**

- Process can't peek into channels and can only wait on one channel at a time
- Output data produced by a process does not depend on the order of its inputs
- Terminates on global deadlock: all process blocked on `wait()`

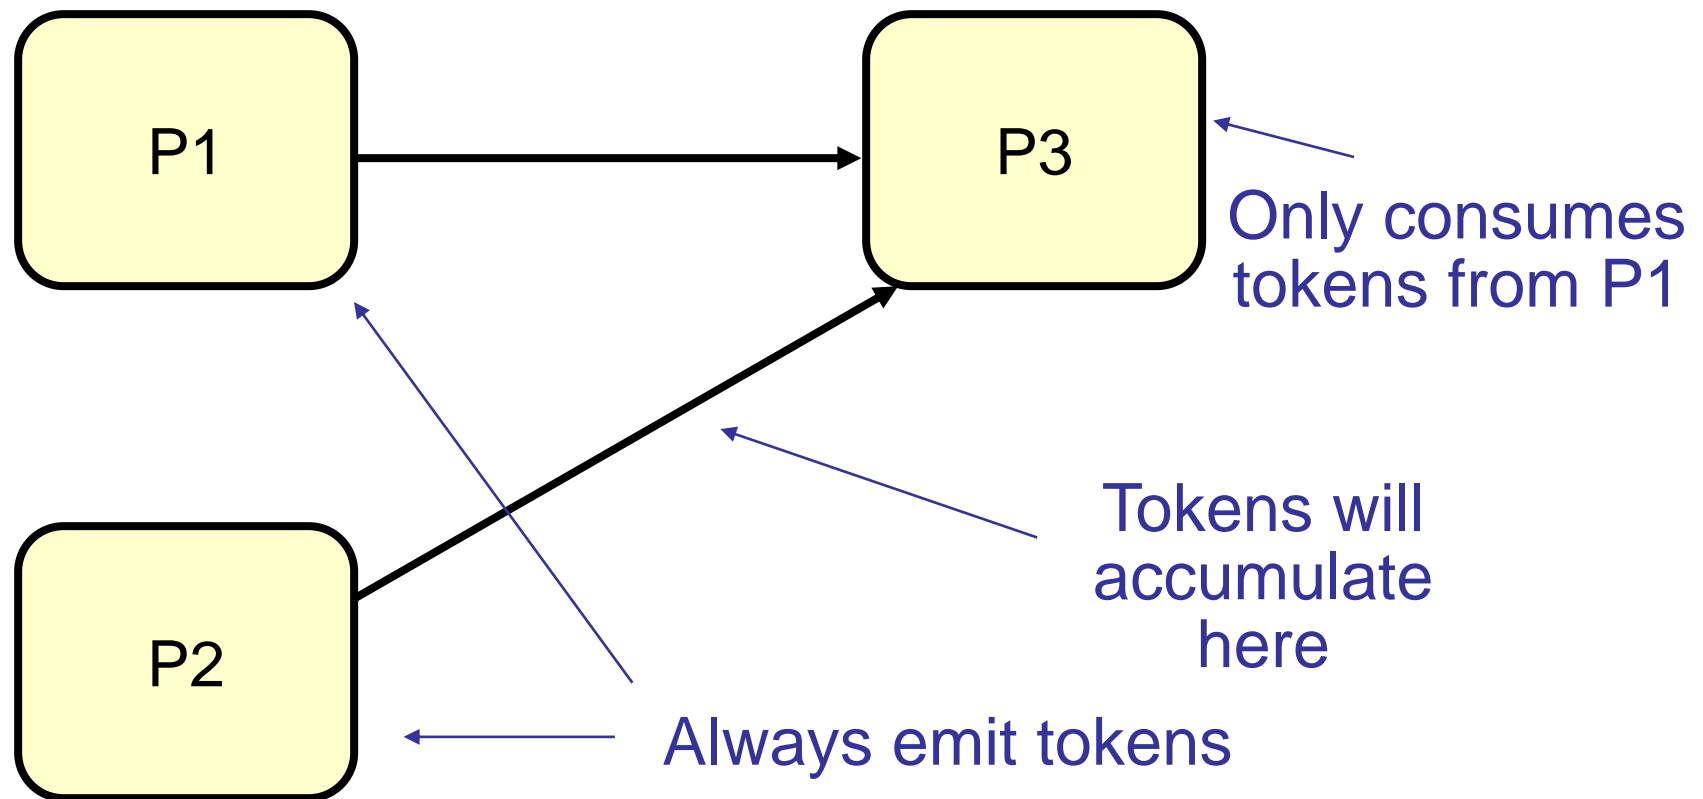
- **Formal mathematical representation**

- Process = continuous function mapping input to output streams

- **Turing-complete, undecidable (in finite time)**

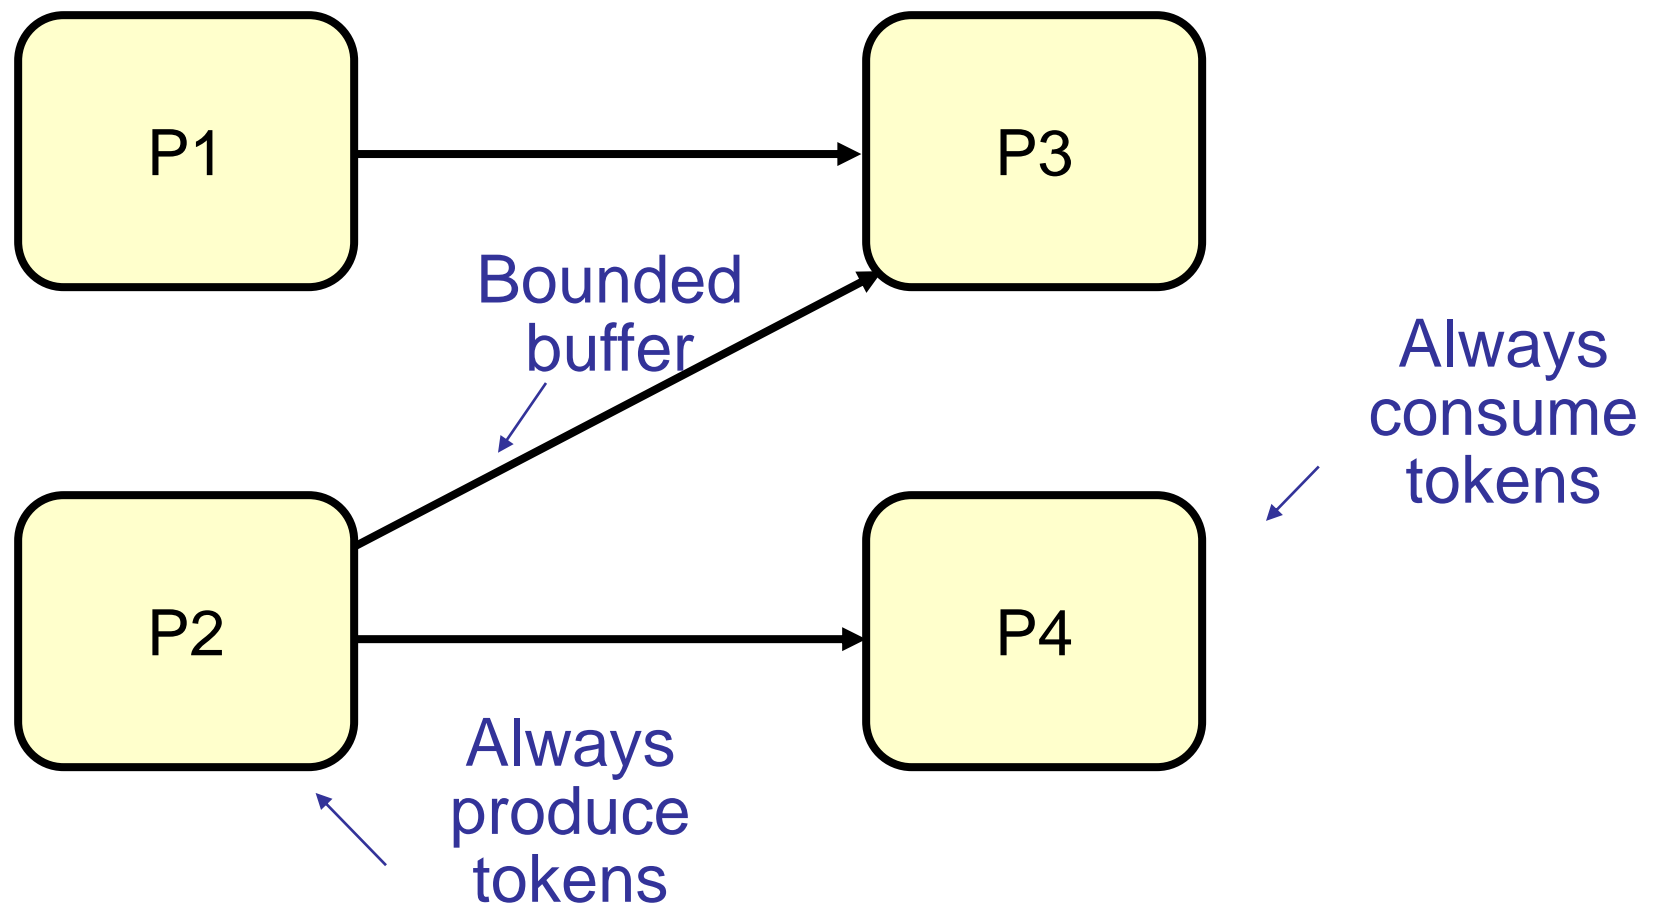
- Terminates?
- Can run in bounded buffers (memory)?

- Scheduling determines memory requirements
- Data-driven scheduling
 - Run processes whenever they are ready:



Source: S. Edwards. "Languages for Digital Embedded Systems", Kluwer 2000.

- **Only run a process whose outputs are being requested**
 - Synchronous, unbuffered message-passing
- **However...**



Source: S. Edwards. "Languages for Digital Embedded Systems", Kluwer 2000.

- **Inherent tradeoffs**


- **Completeness**

- Run processes as long as they are ready
 - Might require unbounded memory

- **Boundedness**


- Block senders when reaching buffer limits
 - Potentially incomplete, artificial deadlocks and early termination
 - Data driven: completeness over boundedness
 - Demand driven: boundedness over completeness and even non-termination

- Start with buffer size 1
- Run P1, P2, P3, P4



Source: S. Edwards. "Languages for Digital Embedded Systems", Kluwer 2000.

- P2 blocked
- Run P1, P3, P1, ... indefinitely



Source: S. Edwards. "Languages for Digital Embedded Systems", Kluwer 2000.

➤ Hybrid approach [Parks95]

- Start with smallest bounded buffers
- Schedule with blocking `send()` until (artificial) deadlock
- Increase size of smallest blocked buffer and continue

• Parks' Scheduling Algorithm (1995)

- Set a capacity on each channel
- Block a write if the channel is full
- Repeat
- Run until deadlock occurs
- If there are no blocking writes → terminate
- Among the channels that block writes, select the channel with least capacity and increase capacity until producer can fire


- **Difficult to implement**

- Size of infinite FIFOs in limited physical memory?
- Dynamic memory allocation, dependent on schedule
- Boundedness vs. completeness vs. non-termination (deadlocks)
- Dynamic context switching

- **Parks' algorithm**

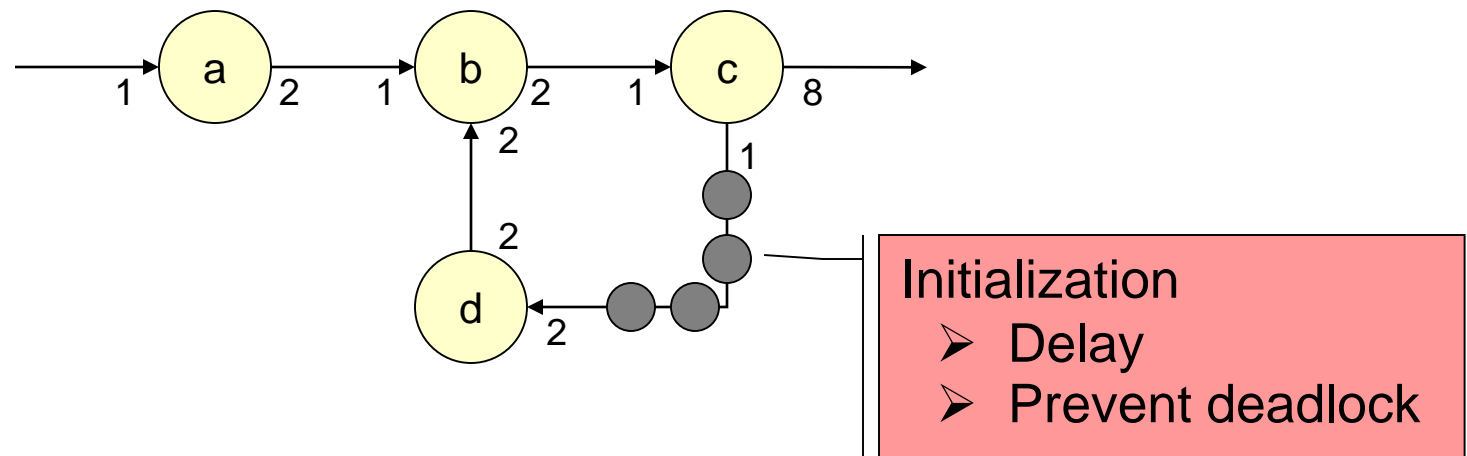
- Non-terminating over bounded over complete execution
 - Does not find every complete, bounded schedule
 - Does not guarantee minimum memory usage
 - Deadlock detection?

- **Breaking processes down into network of *actors***
 - Atomic blocks of computation, executed when *firing*
 - Fire when required number of input *tokens* are available
 - Consume required number of tokens on input(s)
 - Produce number of tokens on output(s)
 - Separate computation & communication/synchronization
 - Actors (indivisible units of computation) may fire simultaneously, any order
 - Tokens (units of communication) can carry arbitrary pieces of data
- **Directed graph of infinite FIFO arcs between actors**
 - Boundedness, completeness, non-termination?



➤ Signal-processing applications

- **Fixed number of tokens per firing**
 - Consume fixed number of inputs
 - Produce fixed number of outputs




- **Can be scheduled statically**
 - Flow of data through system does not depend on values
- **Find a repeated sequence of firings**
 - Run actors in proportion to their rates
 - Fixed buffer sizes, no under- or over-flow

- **Solve system of linear rate equations**

- Balance equations per arc

- $2a = b$
 - $2b = c$
 - $b = d$
 - $2d = c$

➤ $4a = 2b = c = 2d$



- Inconsistent systems

- Only solvable by setting rates to zero
 - Would otherwise (if scheduled dynamically) accumulate tokens

- Underconstrained systems

- Disjoint, independent parts of a design

➤ **Compute repetitions vector**

➤ Linear-time depth-first graph traversal algorithm

- Periodically schedule actors in proportion to their rates

- Smallest integer solution

- $4a = 2b = c = 2d$


- $a = 1, b = 2, c = 4, d = 2$

- Symbolically simulate one iteration of graph until back to initial state

- Insert initialization tokens to avoid deadlock

- $adbccdbcc = a(2db(2c))$

- $a(2db)(4c)$



- Schedule determines memory requirements


- $a(2db(2c))$: 2 token slots on each arc for total of 8 token buffers

- $a(2db)(4c)$: extra initialization tokens, 12 token buffers

- Single appearance schedule to reduce code size

- Looped code generation and compiler optimizations

- **Periodic firings (cyclic pattern of token numbers)**
 - 8:1 Downsampler
 - Traditional SDF: store and consume 8 input tokens for one output




- First firing: consume 1, produce 1
- Second through eighth firing: consume 1, produce 0
- **Static scheduling in similar manner as SDF**


Source: S. Edwards. "Languages for Digital Embedded Systems", Kluwer 2000.

- **Allow actors with boolean control inputs**

- Select actor



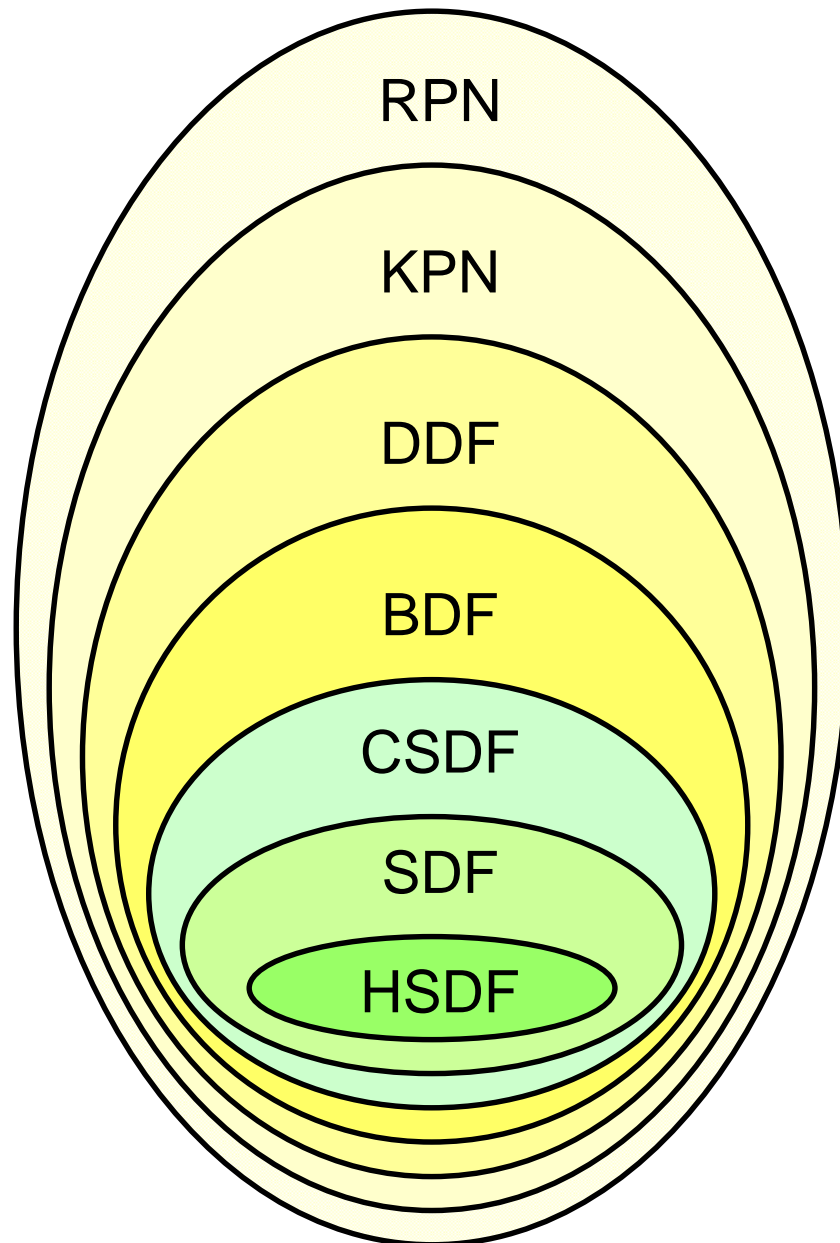
- Switch actor



- **Turing complete**

- Loops, branches
- Quasi-static scheduling

Source: M. Jacome, UT Austin.



Yellow: Turing complete


RPN	Reactive Process Network
KPN	<u>Kahn Process Network</u>
DDF	Dynamic Dataflow
BDF	<u>Boolean Dataflow</u>
CSDF	<u>Cyclo-Static Dataflow</u>
SDF	<u>Synchronous Dataflow</u>
HSDF	Homogeneous SDF (DFG, with single token)

Source: T. Basten, MoCC 2008.

- **Rendezvous-style, synchronous communication**
 - Communicating Sequential Processes (CSP) [Hoare78]
 - Calculus of Communicating Systems (CCS) [Milner80]
 - Restricted interactions
- **Formal, mathematical framework: process algebra**
 - Algebra = <objects, operations, axioms>
 - Objects: processes $\{P, Q, \dots\}$, channels $\{a, b, \dots\}$
 - Composition operators: parallel ($P \parallel Q$), prefix/sequential ($a \rightarrow P$), choice ($P + Q$)
 - Axioms: indemnity ($\emptyset \parallel P = P$), commutativity ($P + Q = Q + P$, $P \parallel Q = Q \parallel P$)
 - Manipulate processes by manipulating expressions
- **Parallel programming languages**
 - CSP-based [Occam/Transputer, Handle-C]

- ✓ Models of Computation (MoCs)
- ✓ Process-based MoCs
- **State-based MoCs**
 - Finite state machines
 - Hierarchical, concurrent state machines
 - Process state machines

- **Status and reactivity (control flow)**
 - **Explicit enumeration of computational states**
 - State represents captured history
 - **Explicit flow of control**
 - Transitions in reaction to events
- **Stepwise operation of a machine**
 - Cycle-by-cycle hardware behavior
 - Finite number of states
 - Not Turing complete
- **State-oriented imperative representation**
 - State only implicit in control/data flow (CDFG)
- **Formal analysis**
 - Reachability, equivalence, ...



- **Finite State Machine (FSM)**


- Basic model for describing control and automata
 - Sequential circuits
- States S , inputs/outputs I/O , and state transitions
 - FSM: $\langle S, I, O, f, h \rangle$
 - Next state function $f: S \times I \rightarrow S$
 - Non-deterministic: f is multi-valued

- **Output function h**

- Mealy-type (input-based), $h: S \times I \rightarrow O$
- Moore-type (state-based), $h: S \rightarrow O$
- Convert Mealy to Moore by splitting states per output

- **Finite State Machine with Data (FSMD)**


- Computation as control and expressions
 - Controller and datapath of RTL processors
- FSM plus variables V
 - FSMD: $\langle S, I, O, \mathbf{V}, f, h \rangle$
 - Next state function $f: S \times V \times I \rightarrow S \times V$
 - Output function $h: S \times V \times I \rightarrow O$



- FSM with Data (FSMD) / Extended FSM (EFSM)**

- Mealy counter


- Implicit self-loops on \bar{e} (absence), $f: S \times V \times 2^I \rightarrow S \times V$, $h: S \times V \times 2^I \rightarrow 2^O$



- Non-Deterministic FSM (NFSM)**

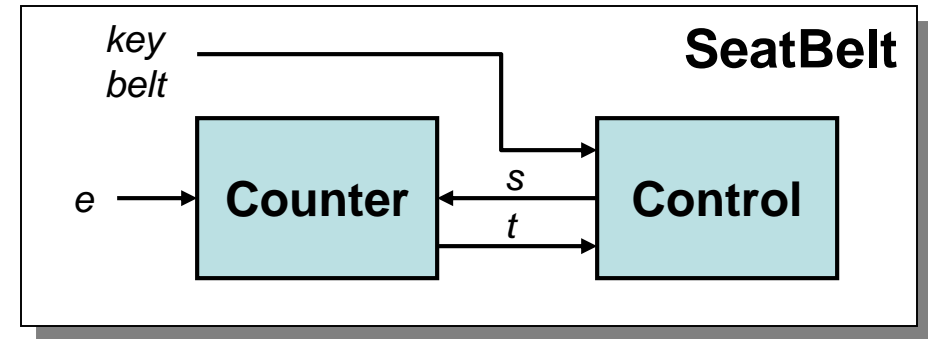
- Choice in control

- Implicit self-loops for unspecified conditions? Usually!
- **Wait:** $belt \ \& \ t$?
- Multiple arcs for same condition?
- Incomplete specification (undecided), unknown behavior (don't care)



• FSM composition

- SeatBelt: $\langle S', I', O', f', h' \rangle$
 - $S' = S_1 \times S_2 = \{ \dots, (\text{Wait}, S1), \dots \}$
 - $I' \subseteq I_1 \cup I_2 = \{e, \text{key}, \text{belt}\}$
 - $O' \subseteq O_1 \cup O_2$
 - $f': S_1 \times S_2 \times I' \rightarrow S_1 \times S_2$, s.t. $f' \in f_1 \times f_2$
 - $h': S_1 \times S_2 \times I' \rightarrow O'$, s.t. $h' \in h_1 \times h_2$



• Connectivity constraints

- Mapping of outputs to inputs: $f_i(s_i, \dots, h_j(s_j, i_j), \dots), h_i(s_i, \dots, h_j(s_j, i_j), \dots)$

➤ Synchronous execution

- Simultaneous and instantaneous, zero delay hypothesis



➤ Composability

➤ Moore

- Well-defined, delayed

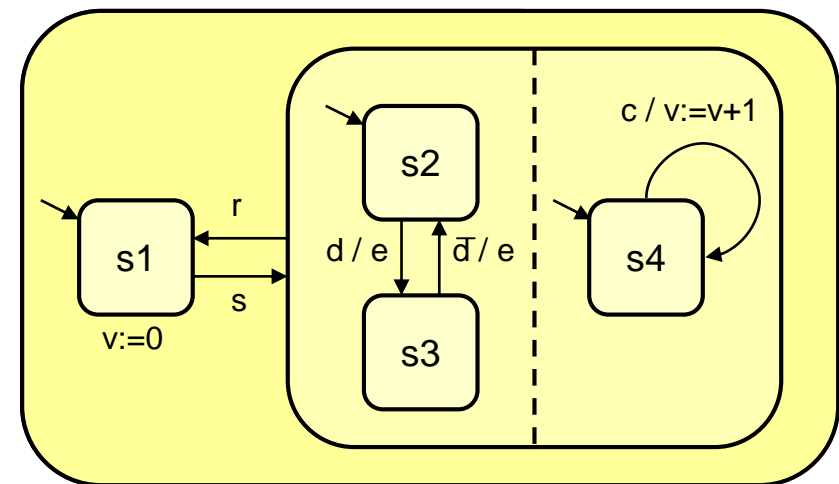
➤ Mealy

- Combinatorial cycles, consistency




Source: M. Jacome, UT Austin.


- **Superstate FSM with Data (SFSMD)**
 - Hierarchy to organize and reduce complexity
 - Superstates that contain complete state machines each
 - Enter into one and exit from any substate
- **Hierarchical Concurrent FSM (HCFSM)**
 - Hierarchical and parallel state composition
 - Lock-step concurrent composition and execution
 - Communication through global variables, signals and events
- Graphical notation
[StateCharts]




- Hierarchy (OR state)**



- Concurrency (AND state)**



- **Reaction time upon external event?**
 - Synchronous, reactive: zero time, event broadcast



- **Event propagation?**
 - Grandfather paradox
 - Inconsistency or non-determinism
 - Synchronous
 - Reject cycles [Argos, Lustre]
 - Require fixed-point [SyncCharts/Esterel]

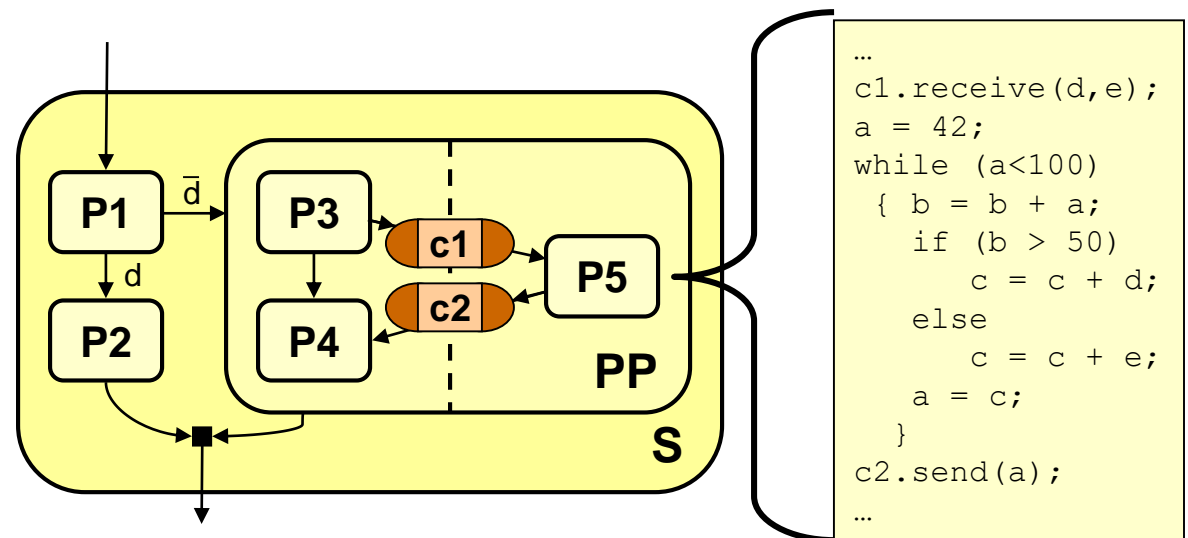
- ***N* micro-steps (internal) per macro-step (I/O) [Statemate]**
 - Events posted in next and only in next micro step
 - “Synchronous”
 - One micro/macro step at regular times: delayed reaction, *not* synchronous (Moore)
 - “Asynchronous”
 - Zero-delay micro steps: causal chain reaction (Mealy, but: state updates in cycles)
 - Deterministic
 - Together with other rules, e.g. priority of conflicting transitions

Source: “Statemate Course,” K. Baukus

- **Sound combination of process and state based models**
 - Asynchronous concurrent HCFSM execution [UML]
 - Explicit event queues, deadlock analysis [PetriNet]
 - Globally asynchronous, locally synchronous (GALS)
 - Co-design Finite State Machines (CFSM) [Polis]
 - Leaf states are imperative processes
 - Program State Machine (PSM) [SpecSyn]
 - Transition-Immediately (TI) and Transition-on-Completion (TOC)

➤ Processes and abstract channels

- Computation & communication
- Process State Machine (PSM) [SpecC]



- **Programming models**
 - Imperative, declarative: C, C++, Haskell, ...
 - Synchronous, reactive: Esterel, Lustre
 - Asynchronous threads: Java
- **Process-based models: KPN, Dataflow, SDF**
 - Data dominated, block diagram level
- **State-based models: FSM(D), HCFSM**
 - Control dominated, machine level
- **Hybrid models**
 - Combination of process and state (data and control)
 - Behavior from specification down to implementation